

```

1 import torch
2 import torch.nn as nn
3 import torch.optim as optim
4 from torch.utils.data import Dataset, DataLoader
5 from collections import defaultdict
6 import random
7
8
9 def read_tsv(path):
10     data = []
11     with open(path, encoding='utf-8') as f:
12         for line in f:
13             dev, lat, freq = line.strip().split('\t')
14             data.extend([(lat, dev)] * int(freq))
15     return data
16
17 def build_vocab(sequences):
18     vocab = {'<pad>': 0, '<sos>': 1, '<eos>': 2}
19     for seq in sequences:
20         for char in seq:
21             if char not in vocab:
22                 vocab[char] = len(vocab)
23     return vocab
24
25
26 class TransliterationDataset(Dataset):
27     def __init__(self, data, input_vocab, target_vocab):
28         self.data = data
29         self.input_vocab = input_vocab
30         self.target_vocab = target_vocab
31
32     def __len__(self):
33         return len(self.data)
34
35     def encode_seq(self, seq, vocab, add_sos_eos=False):
36         ids = [vocab[c] for c in seq]
37         if add_sos_eos:
38             ids = [vocab['<sos>']] + ids + [vocab['<eos>']]
39         return torch.tensor(ids, dtype=torch.long)
40
41     def __getitem__(self, idx):
42         latin, dev = self.data[idx]
43         return self.encode_seq(latin, self.input_vocab), self.encode_seq(dev, self.target_vocab, True)
44
45 def collate_fn(batch):
46     srcs, trgs = zip(*batch)
47     srcs_padded = nn.utils.rnn.pad_sequence(srcs, batch_first=True, padding_value=0)
48     trgs_padded = nn.utils.rnn.pad_sequence(trgs, batch_first=True, padding_value=0)
49     return srcs_padded, trgs_padded
50
51
52 class Encoder(nn.Module):
53     def __init__(self, input_dim, emb_dim, hid_dim, n_layers, rnn_type='gru'):
54         super().__init__()
55         self.embedding = nn.Embedding(input_dim, emb_dim)
56         rnn_cls = {'rnn': nn.RNN, 'lstm': nn.LSTM, 'gru': nn.GRU}[rnn_type]
57         self.rnn = rnn_cls(emb_dim, hid_dim, n_layers, batch_first=True)
58         self.rnn_type = rnn_type
59
60     def forward(self, src):
61         embedded = self.embedding(src)
62         outputs, hidden = self.rnn(embedded)
63         return hidden
64
65 class Decoder(nn.Module):
66     def __init__(self, output_dim, emb_dim, hid_dim, n_layers, rnn_type='gru'):
67         super().__init__()
68         self.embedding = nn.Embedding(output_dim, emb_dim)
69         rnn_cls = {'rnn': nn.RNN, 'lstm': nn.LSTM, 'gru': nn.GRU}[rnn_type]
70         self.rnn = rnn_cls(emb_dim, hid_dim, n_layers, batch_first=True)
71         self.fc_out = nn.Linear(hid_dim, output_dim)
72
73     def forward(self, input, hidden):
74         input = input.unsqueeze(1)
75         embedded = self.embedding(input)
76         output, hidden = self.rnn(embedded, hidden)
77         prediction = self.fc_out(output.squeeze(1))
78         return prediction, hidden
79
80 class Seq2Seq(nn.Module):
81     def __init__(self, encoder, decoder, device):
82         super().__init__()
83         self.encoder = encoder
84         self.decoder = decoder

```

```

85     self.device = device
86
87     def forward(self, src, trg, teacher_forcing_ratio=0.5):
88         batch_size, trg_len = trg.shape
89         output_dim = self.decoder.fc_out.out_features
90         outputs = torch.zeros(batch_size, trg_len, output_dim).to(self.device)
91
92         hidden = self.encoder(src)
93         input = trg[:, 0]
94
95         for t in range(1, trg_len):
96             output, hidden = self.decoder(input, hidden)
97             outputs[:, t] = output
98             teacher_force = random.random() < teacher_forcing_ratio
99             top1 = output.argmax(1)
100            input = trg[:, t] if teacher_force else top1
101
102        return outputs
103
104
105    def train(model, data_loader, optimizer, criterion, clip=1):
106        model.train()
107        epoch_loss = 0
108        for src, trg in data_loader:
109            src, trg = src.to(model.device), trg.to(model.device)
110            optimizer.zero_grad()
111            output = model(src, trg)
112            output_dim = output.shape[-1]
113            output = output[:, 1:].reshape(-1, output_dim)
114            trg = trg[:, 1:].reshape(-1)
115            loss = criterion(output, trg)
116            loss.backward()
117            torch.nn.utils.clip_grad_norm_(model.parameters(), clip)
118            optimizer.step()
119            epoch_loss += loss.item()
120        return epoch_loss / len(data_loader)
121
122    def accuracy(model, data_loader):
123        model.eval()
124        correct, total = 0, 0
125        with torch.no_grad():
126            for src, trg in data_loader:
127                src, trg = src.to(model.device), trg.to(model.device)
128                output = model(src, trg, 0)
129                preds = output.argmax(-1)
130                for pred, true in zip(preds, trg):
131                    if torch.equal(pred[1:], true[1:]):
132                        correct += 1
133                total += 1
134        return correct / total
135
136    def predict(model, src_seq, input_vocab, output_vocab, max_len=30):
137        model.eval()
138        inv_vocab = {v: k for k, v in output_vocab.items()}
139        src_tensor = torch.tensor([input_vocab[c] for c in src_seq], dtype=torch.long).unsqueeze(0).to(model.device)
140        hidden = model.encoder(src_tensor)
141        input = torch.tensor([output_vocab['<sos>']], device=model.device)
142        output = []
143        for _ in range(max_len):
144            out, hidden = model.decoder(input, hidden)
145            top1 = out.argmax(1)
146            char = inv_vocab[top1.item()]
147            if char == '<eos>':
148                break
149            output.append(char)
150            input = top1
151        return ''.join(output)
152
153
154    DEVICE = torch.device("cuda" if torch.cuda.is_available() else "cpu")
155
156    train_path = "/content/hi.translit.sampled.train.tsv"
157    val_path = "/content/hi.translit.sampled.dev.tsv"
158    train_data = read_tsv(train_path)
159    val_data = read_tsv(val_path)
160
161    input_vocab = build_vocab([d[0] for d in train_data])
162    target_vocab = build_vocab([d[1] for d in train_data])
163
164    train_dataset = TransliterationDataset(train_data, input_vocab, target_vocab)
165    val_dataset = TransliterationDataset(val_data, input_vocab, target_vocab)
166    train_loader = DataLoader(train_dataset, batch_size=32, shuffle=True, collate_fn=collate_fn)
167    val_loader = DataLoader(val_dataset, batch_size=32, shuffle=False, collate_fn=collate_fn)
168
169    INPUT_DIM = len(input_vocab)
170    OUTPUT_DIM = len(target_vocab)

```

```

171 EMB_DIM = 64
172 HID_DIM = 128
173 N_LAYERS = 1
174 RNN_TYPE = 'gru'
175
176 encoder = Encoder(INPUT_DIM, EMB_DIM, HID_DIM, N_LAYERS, RNN_TYPE)
177 decoder = Decoder(OUTPUT_DIM, EMB_DIM, HID_DIM, N_LAYERS, RNN_TYPE)
178 model = Seq2Seq(encoder, decoder, DEVICE).to(DEVICE)
179
180 optimizer = optim.Adam(model.parameters(), lr=0.001)
181 criterion = nn.CrossEntropyLoss(ignore_index=target_vocab['<pad>'])
182
183 for epoch in range(10):
184     loss = train(model, train_loader, optimizer, criterion)
185     acc = accuracy(model, val_loader)
186     print(f"Epoch {epoch+1} | Loss: {loss:.4f} | Val Accuracy: {acc:.4f}")
187
188 print("\nSample Predictions:")
189 for i in range(5):
190     src_sample, tgt_sample = val_data[i]
191     pred = predict(model, src_sample, input_vocab, target_vocab)
192     print(f"Input: {src_sample} | Target: {tgt_sample} | Predicted: {pred}")
193

```

```

🔄 Epoch 1 | Loss: 1.5929 | Val Accuracy: 0.0069
Epoch 2 | Loss: 0.8821 | Val Accuracy: 0.0103
Epoch 3 | Loss: 0.7274 | Val Accuracy: 0.0162
Epoch 4 | Loss: 0.6378 | Val Accuracy: 0.0167
Epoch 5 | Loss: 0.5878 | Val Accuracy: 0.0204
Epoch 6 | Loss: 0.5420 | Val Accuracy: 0.0238
Epoch 7 | Loss: 0.5066 | Val Accuracy: 0.0199
Epoch 8 | Loss: 0.4810 | Val Accuracy: 0.0216
Epoch 9 | Loss: 0.4545 | Val Accuracy: 0.0219
Epoch 10 | Loss: 0.4318 | Val Accuracy: 0.0232

```

Sample Predictions:

```

Input: ankan | Target: अंकन | Predicted: अंकन
Input: ankan | Target: अंकन | Predicted: अंकन
Input: ankan | Target: अंकन | Predicted: अंकन
Input: angkor | Target: अंगकोर | Predicted: अंगकर
Input: angkor | Target: अंगकोर | Predicted: अंगकर

```

```

1 import pandas as pd
2 import re
3 from datasets import load_dataset
4 from transformers import GPT2Tokenizer, GPT2LMHeadModel, Trainer, TrainingArguments, pipeline
5
6
7 def load_and_clean_lyrics(*csv_paths):
8     dfs = [pd.read_csv(path) for path in csv_paths]
9     lyrics_df = pd.concat(dfs)
10
11     def clean_lyrics(lyric):
12         if pd.isna(lyric):
13             return ""
14         lyric = str(lyric)
15         lyric = re.sub(r'^#+', '', lyric)
16         lyric = lyric.encode('utf-8').decode('utf-8', 'ignore')
17         lyric = re.sub(r'[\u2018\u2019\u201c\u201d]+', '', lyric)
18         lyric = re.sub(r'^\x00-\x7F+', '', lyric)
19         return lyric.strip()
20
21     return lyrics_df['Lyric'].dropna().apply(clean_lyrics).tolist()
22
23 lyrics_texts = load_and_clean_lyrics('/content/EdSheeran.csv', '/content/JustinBieber.csv')
24
25 with open("lyrics_dataset.txt", "w", encoding="utf-8") as f:
26     for lyric in lyrics_texts:
27         f.write(lyric + "\n\n")
28
29 dataset = load_dataset("text", data_files={"train": "lyrics_dataset.txt"})
30
31 tokenizer = GPT2Tokenizer.from_pretrained("gpt2")
32 tokenizer.pad_token = tokenizer.eos_token
33
34 def tokenize_function(example):
35     return tokenizer(example["text"], truncation=True, padding="max_length", max_length=512)
36
37 tokenized_dataset = dataset.map(tokenize_function, batched=True, remove_columns=["text"])
38 tokenized_dataset = tokenized_dataset.map(lambda examples: {"labels": examples["input_ids"]}, batched=True)
39
40 model = GPT2LMHeadModel.from_pretrained("gpt2")
41
42 training_args = TrainingArguments(
43     output_dir="./gpt2-lyrics",
44     per_device_train_batch_size=2,

```

```
45     num_train_epochs=3,
46     logging_steps=100,
47     save_steps=500,
48     save_total_limit=1,
49     prediction_loss_only=True,
50     report_to="none",
51     fp16=False
52 )
53
54 trainer = Trainer(
55     model=model,
56     args=training_args,
57     train_dataset=tokenized_dataset["train"],
58     tokenizer=tokenizer
59 )
60
61 trainer.train()
62
63
64
65 generator = pipeline("text-generation", model=model, tokenizer=tokenizer)
66 output = generator("I remember those nights when", max_length=100, num_return_sequences=1)[0]["generated_text"]
67 print(output)
68
```

Generating train split: 1282/0 [00:00<00:00, 34168.72 examples/s]

/usr/local/lib/python3.11/dist-packages/huggingface\_hub/utils/\_auth.py:94: UserWarning:

The secret `HF\_TOKEN` does not exist in your Colab secrets.

To authenticate with the Hugging Face Hub, create a token in your settings tab (<https://huggingface.co/settings/tokens>), set it as secret in your Colab secrets. You will be able to reuse this secret in all of your notebooks.

Please note that authentication is recommended but still optional to access public models or datasets.

warnings.warn(

tokenizer\_config.json: 100% 26.0/26.0 [00:00<00:00, 3.00kB/s]

vocab.json: 100% 1.04M/1.04M [00:00<00:00, 5.18MB/s]

merges.txt: 100% 456k/456k [00:00<00:00, 3.18MB/s]

tokenizer.json: 100% 1.36M/1.36M [00:00<00:00, 9.08MB/s]

config.json: 100% 665/665 [00:00<00:00, 80.3kB/s]

Map: 100% 1282/1282 [00:01<00:00, 714.99 examples/s]

Map: 100% 1282/1282 [00:00<00:00, 2836.24 examples/s]

Xet Storage is enabled for this repo, but the 'hf\_xet' package is not installed. Falling back to regular HTTP download. For better performance, WARNING:huggingface\_hub.file\_download:Xet Storage is enabled for this repo, but the 'hf\_xet' package is not installed. Falling back to regular

model.safetensors: 100% 548M/548M [00:01<00:00, 310MB/s]

generation\_config.json: 100% 124/124 [00:00<00:00, 8.76kB/s]

<ipython-input-5-c9ee5be06c98>:70: FutureWarning: `tokenizer` is deprecated and will be removed in version 5.0.0 for `Trainer.\_\_init\_\_`. Use `p  
trainer = Trainer(